

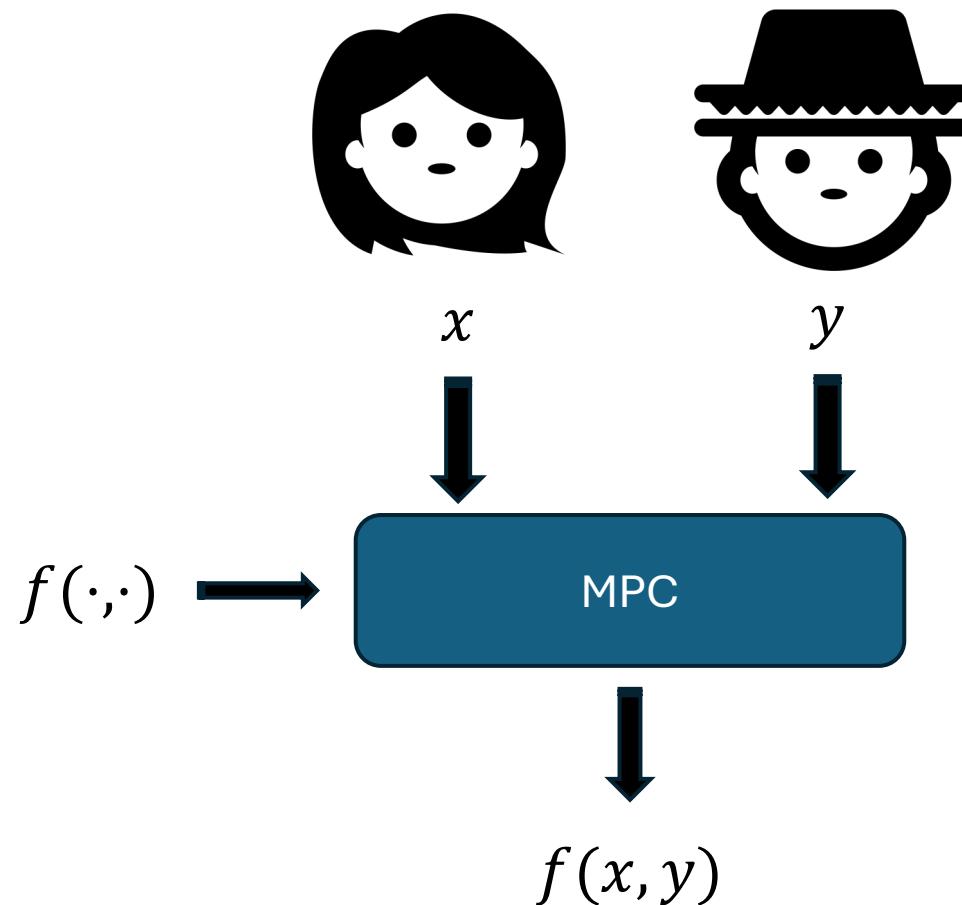
# Toss: Garbled PIR from Table-Only Stacking

**Lucien K. L. Ng** and Vladimir Kolesnikov

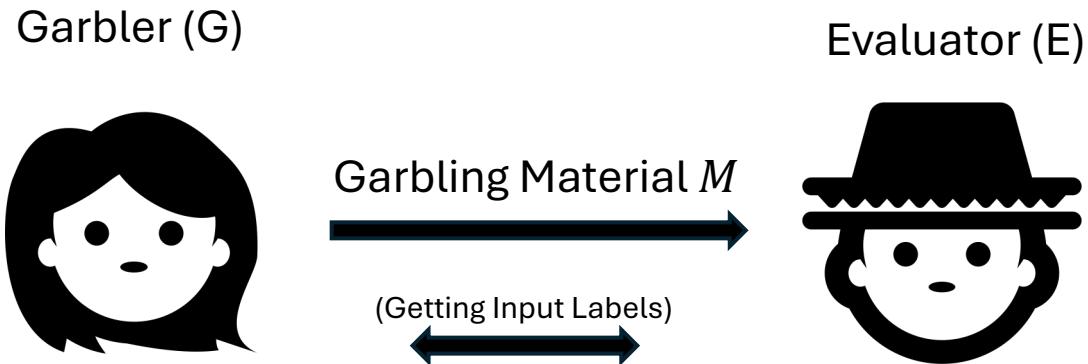
Georgia Institute of Technology



# Multi-Party Computation (MPC)



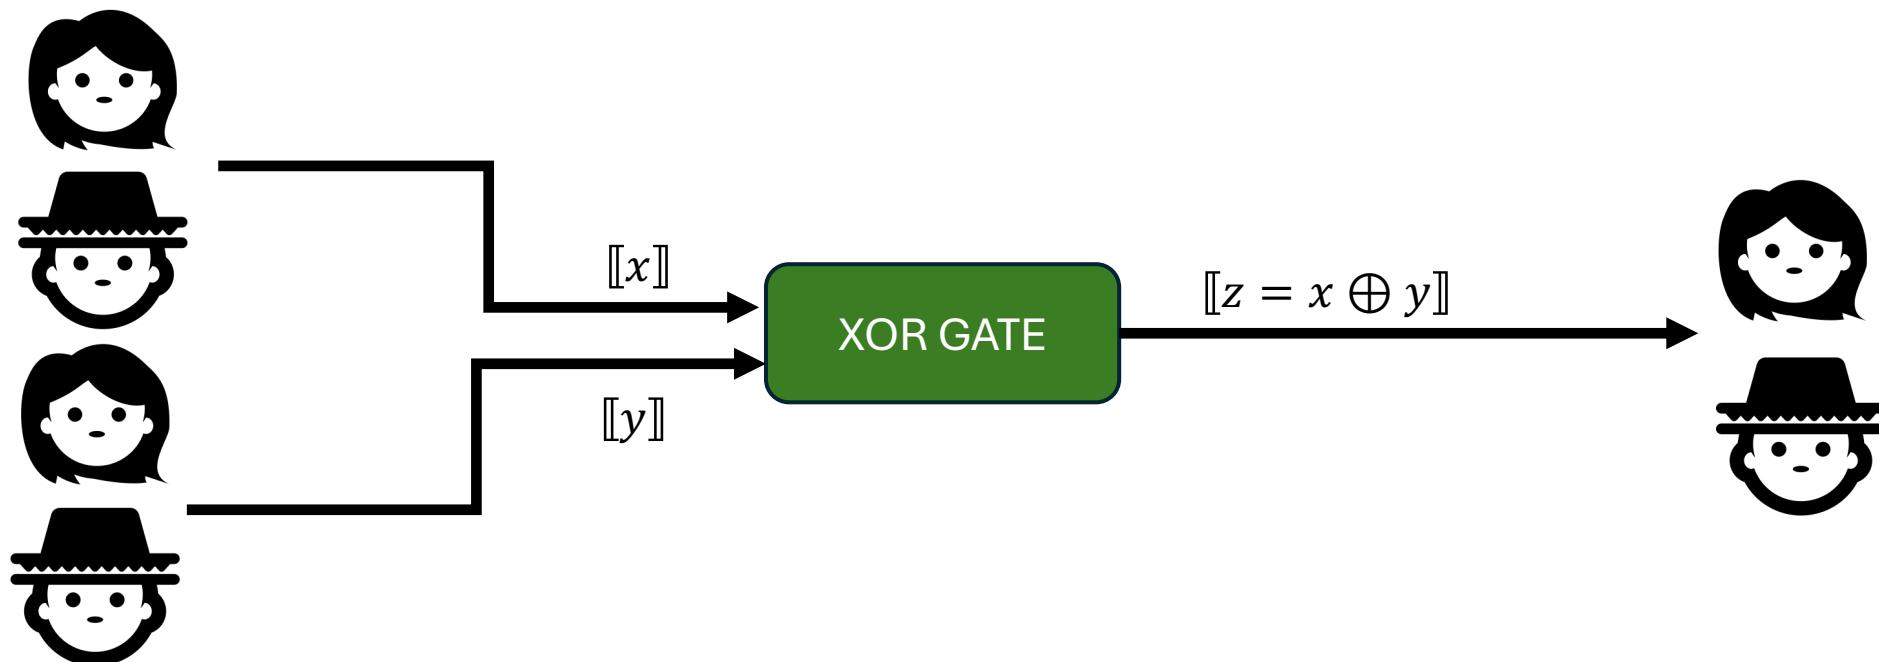
# Garbled Circuit (GC)



- Constant Rounds
- Symmetric Key Primitives
- Authenticity
- Limited Choice of Gates

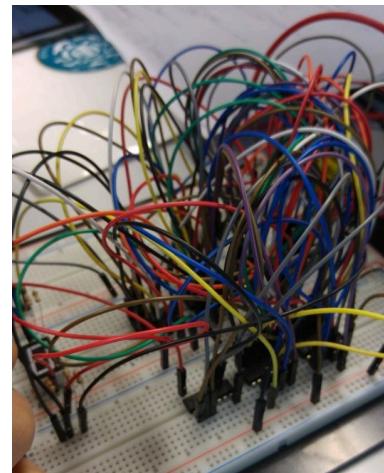
# Circuits as a Composition of Gates

- XOR gates:  $XOR(x, y) = x \oplus y$
- AND gates:  $AND(x, y) = x \cdot y$
- $x, y \in \{0, 1\}$

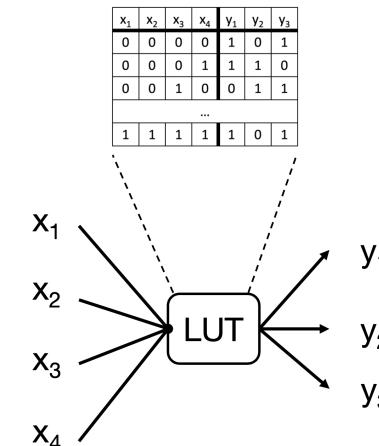


# How about Garbled Lookup Table (LUT)?

- Parameters: Lookup Table  $F$  of size  $N = 2^n$
- Input:  $G$  and  $E$  hold garbled share index  $\llbracket x \rrbracket$ 
  - $x$  may have  $n \geq 1$  bit
- Output:  $G$  and  $E$  receive  $\llbracket F[x] \rrbracket$



VS.



# GLUT Applications

- Privacy-Preserving Machine Learning (PPML)
  - LUTs are used to compute highly-complicated functions and operations
  - SiRnn [SP'21] uses 1020-row LUTs for  $1/\sqrt{x}$
  - Beacon [USS'23] uses  $2^{12}$ -row LUTs for tanh and sigmoid over half-floats
  - Sigma [PETS'24] uses  $2^{13}$ - and  $2^{16}$ -row LUTs for  $1/x$
- Client-Malicious PPML due to GC's authenticity
  - MUSE [USS'21], SIMC [USS'22]
- Privacy-Preserving Deterministic Finite Automata
  - Substring Search (KMP Algorithm)
  - DNA Pattern Matching
- New Opportunity to Optimize Circuits

# Prior Art of Garbled LUT/PIR

- For now, we focus on communication cost
- Naive (Yao's [FOCS'86]):  $2^n m \kappa$  bits
  - logrow's [EC'24] :  $\approx nm\kappa + 2^n m$  bits
    - Security Parameter  $\kappa$  is usually 128

<sup>no  $\kappa$</sup>   
but as large as the table

$n$  = # input bits  
 $m$  = # output bits

*Can we reduce the communication cost to be  
sublinear in table size and  $\kappa$ ?*

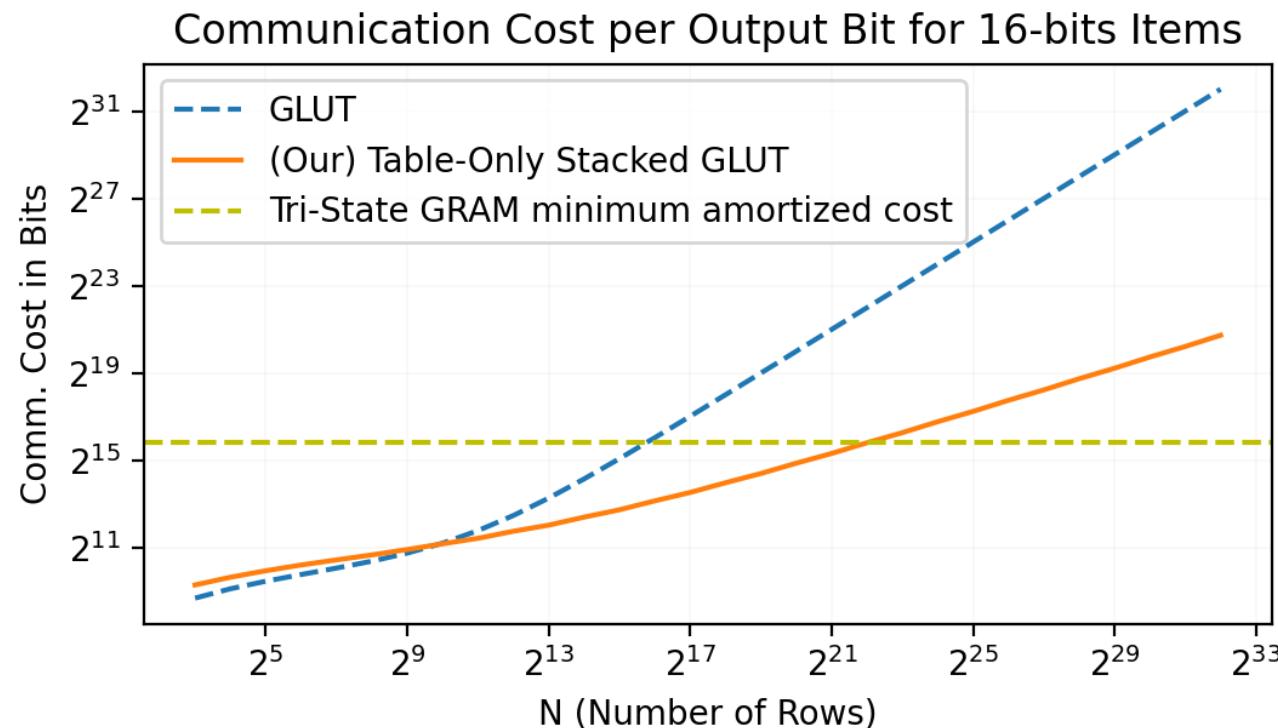
# Recent Efforts to Avoid $\kappa$

Scheme	Optimization Target	Crypto Assumption
<b>Ours</b>	LUT	😊 (Symmetric) CCRH
BitGC [EC'25]	AND	(Asymmetric) RLWE/NTRU
Silent Circuit Relinearisation [Crypto'25]	AND/Mult/LUT	(Asymmetric) DCR
Breaking the $1/\lambda$ -Rate Barrier for Arithmetic Garbling [EC'25]	Mult	(Asymmetric) Power-DDH
$\omega(1/\lambda)$ -Rate Boolean Garbling Scheme from Generic Groups [EC'25]	Mult	(Asymmetric) Generic Group model
Rate-1 Arithmetic Garbling From Homomorphic Secret Sharing [TCC'24]	Mult	(Asymmetric) DCR
New Ways to Garble Arithmetic Circuits [EC'25]	Mult	(Asymmetric) DCR/LWE

- XOR or Addition are ignored as they are communication-free
- Mult: Arithmetic Multiplication Gate
- AND: Boolean AND Gate

# Toss' (Ours) Communication Cost

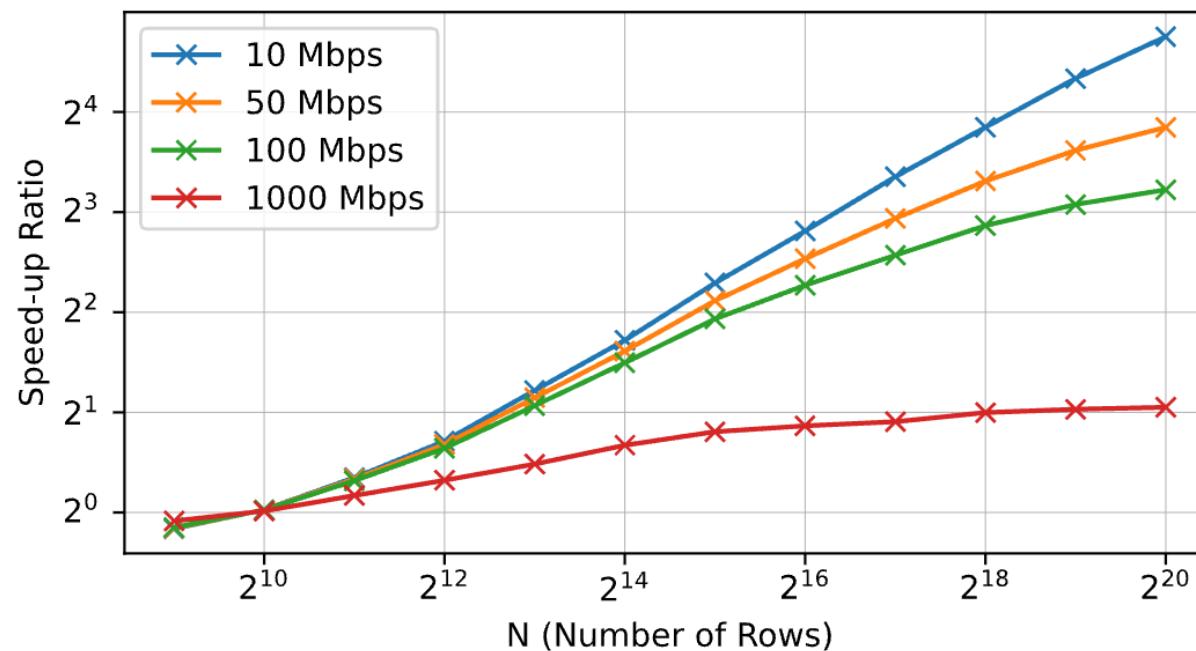
- Toss reduces the comm. cost to  $O(\sqrt{N}m\sqrt{\kappa})$  bits
  - with a lean hidden constant (< 10)



- Assume
- $m = 16$  bits items
  - $\kappa = 128$

# Concrete Runtime Improvement

- Throughput increases from  $\approx 10.6$  lookup/s to  $\approx 81$  lookup/s ( $> 7.5$ x)
  - For  $N = 2^{20}$  with Laptop (M3 Pro Macbook) + 100Mbps Network



Assume

- $m = 16$  bits items
- $\kappa = 128$

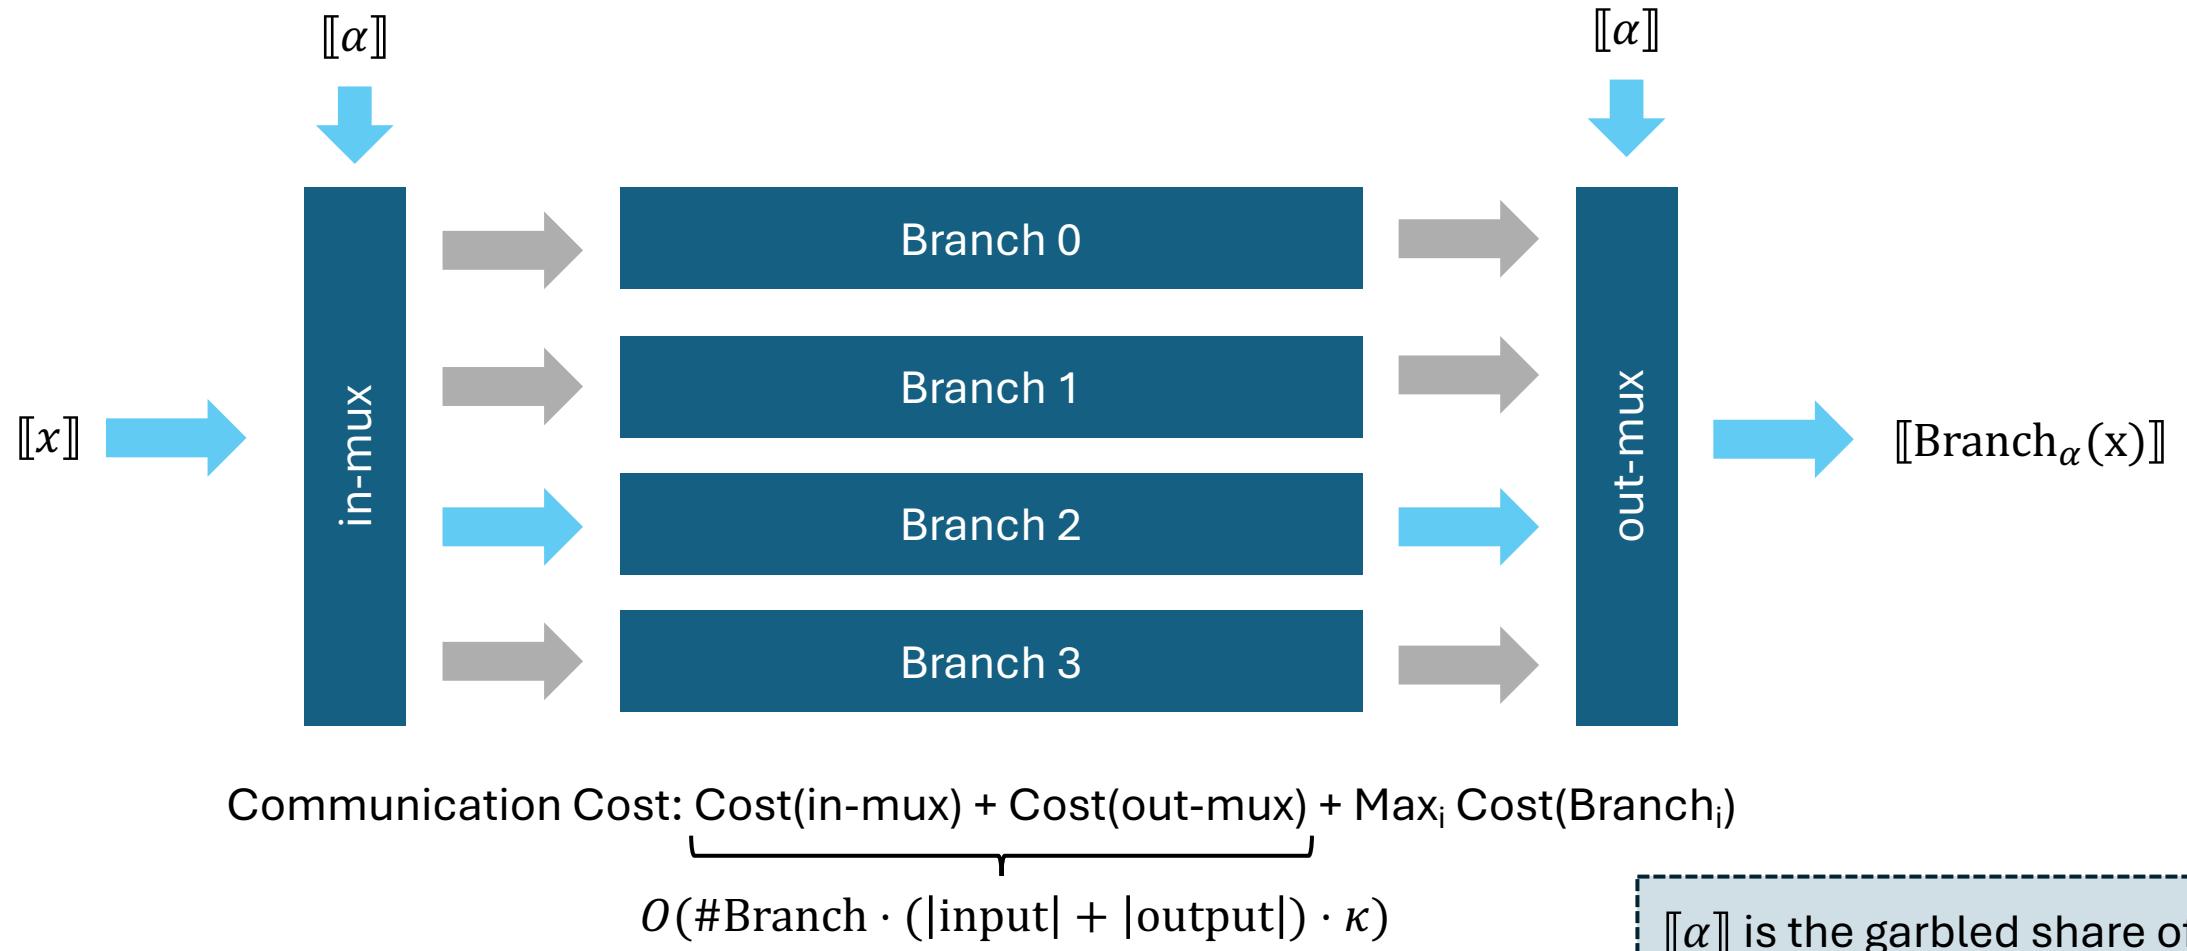
# Garbled PIR vs. LUT

- GPIR : Garbled Private Information Retrieval

	GPIR	GLUT
Table Content	Known by both G & E (i.e., <b>Public</b> )	Known by G
Table Size	<b>Large</b> (e.g., $N > 2^{12}$ )	Small
Scheme	<b>Toss (Ours)</b>	logrow Yao's

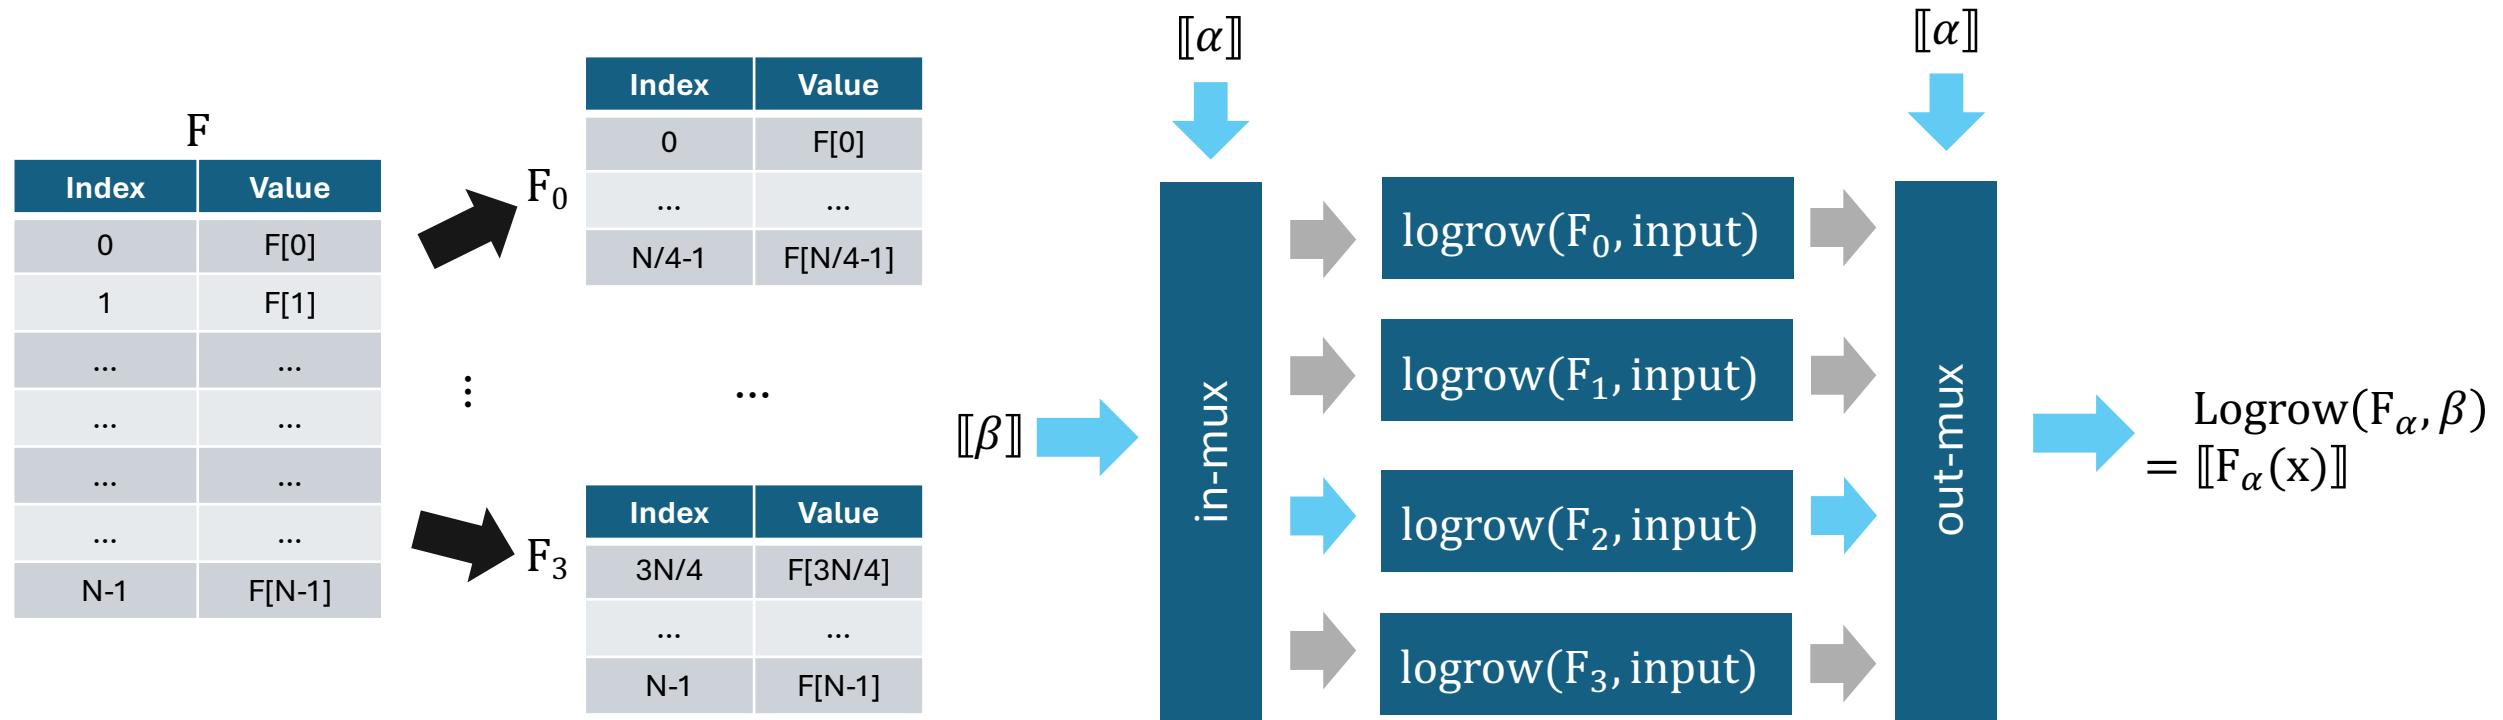
# Ingredient: Stacked Garbling [Crypto'20]

Goal: Reduce communication cost for conditional programs



# Basic Idea: Stacking logrow

- Decompose the lookup index  $\llbracket x \rrbracket = \llbracket \alpha \rrbracket \parallel \llbracket \beta \rrbracket$



# Communication Cost of Stacked logrow

$$\text{Cost(in-mux)} + \text{Cost(out-mux)} + \max_i \text{Cost(Branch}_i\text{)}$$

$O(B \cdot m \cdot \kappa)$

$\text{Cost}(\text{logrow}(F_i, \cdot))$   
 $= \log \frac{N}{B} m \kappa + \frac{N}{B} m$

Plugging in  $B = \sqrt{N/\kappa}$   $\rightarrow$  Cost =  $O(\sqrt{N}m\sqrt{\kappa})$

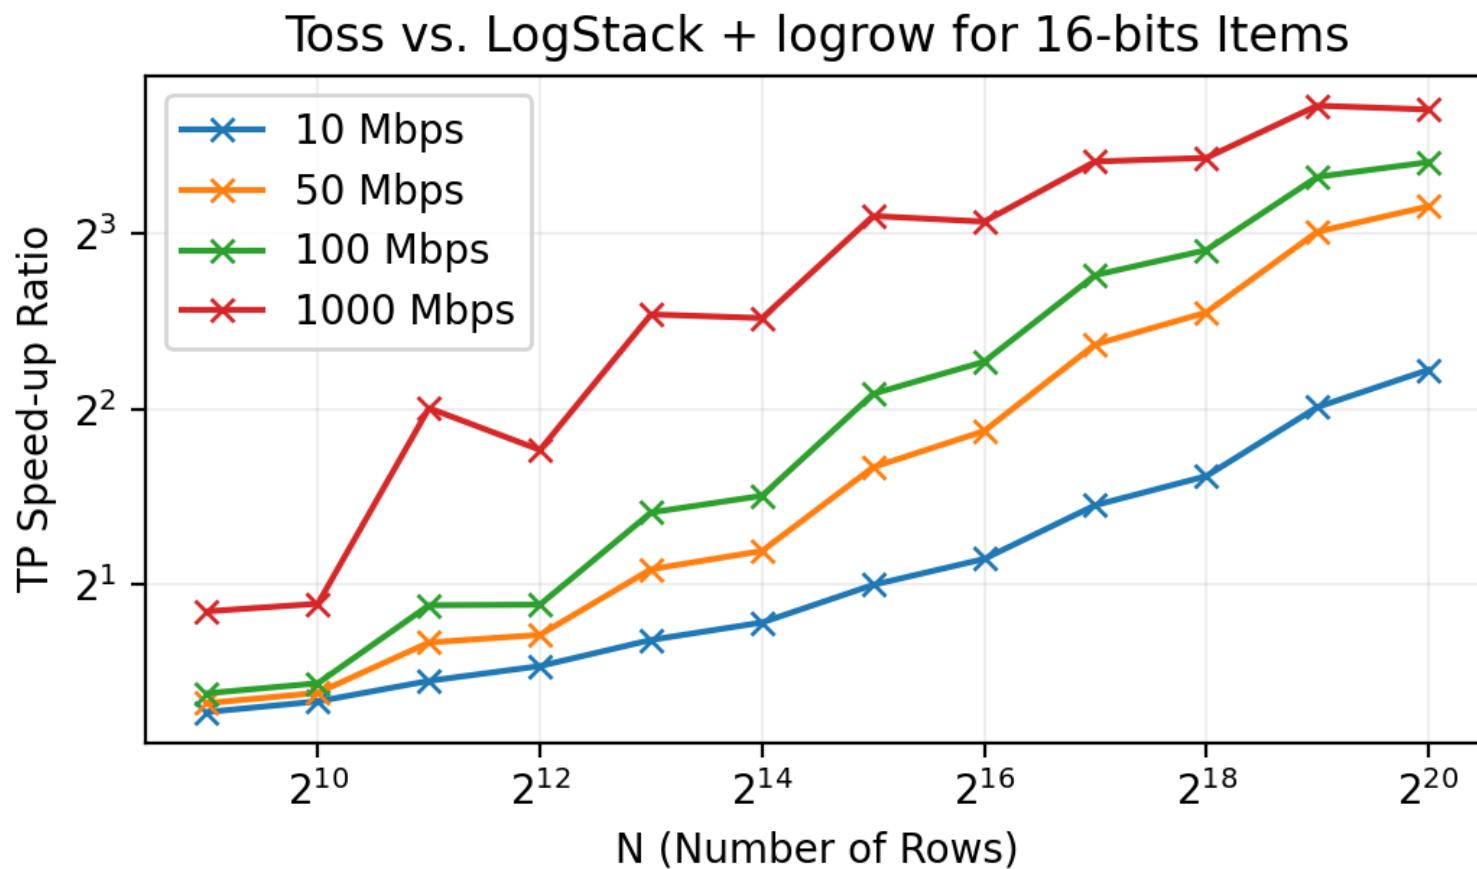
- $B$  = # Branch/Sub-table
- $m$  = item size
- $N = 2^n$  = table size

# ...What about Computation?

- Original Stacked Garbling [Crypto'20]:  $O(B^2 \cdot |\text{Branch}|)$
- LogStack [EC'21]:  $O(B \log B \cdot |\text{Branch}|)$
- For our stacked logrow, the comp. cost becomes  $O(\log \frac{N}{\kappa} \cdot Nm\kappa)$ 
  - vs. unstacked logrow's  $O(Nm\kappa)$
  - E.g., for  $N = 2^{24}$ , it is stacked logrow is 8.5x slower.

*Can we avoid the  $O\left(\log \frac{N}{\kappa}\right)$  computation overhead?*

# Stacked logrow vs. TOSS



# Toss: Table-only Stacking of Sub-tables

- We stack only the sub-tables instead of the entire logrow
  - Much more light-weight!
  - Enabling more low-level algorithmic optimization
- Toss still has  $O(\log B)$  overhead
  - But it's on the plain sub-tables instead of on sub-circuit with  $\kappa$ -blowup
  - Toss's Stacking:  $O(\log B) \cdot O(Nm)$  vs. Naïve:  $O(\log B) \cdot O(Nm\kappa)$
- (Still, the comp. cost is  $O(Nm\kappa)$  due to multiplexing gadgets)

# Closing Remark of Toss

- What else in our paper
  - Detailed ideas of table-only stacking
  - How to reduce the cost for multiplexing in stacked garbling
- Summary:
  - Toss reduces communication without adding more computation
  - Part of the effort of avoid  $\kappa$ -factor in GC communication cost
    - Toss uses only symmetric-key primitive!
  - We use stacked garbling as an important building block
- Ad: Another paper @ Blockchain and Distributed Systems #2
  - Lite-PoT: Practical Powers-of-Tau Setup Ceremony